



NLP MIMIC III

MAZIYAR MIRZAEI

AI IN HEALTHCARE

FALL 2024 UT AUSTIN

ANALYZING DIABETES PATIENT NOTES WITH NLP TECHNIQUES

- **Objective:** To process clinical notes of diabetes patients from the MIMIC-III dataset using NLP methods such as SpaCy, SciSpaCy, Word2Vec, and BERT models for entity recognition and embedding visualization.

DATA LOADING

- Loading Clinical Notes and Diagnoses Data
- **Explanation:** Load the first 50,000 rows of the NOTEEVENTS.csv file and the entire DIAGNOSES_ICD.csv file into pandas DataFrames for processing.

```
import pandas as pd

# Load the data
data = pd.read_csv('NOTEEVENTS.csv', nrows=50000)

# Load diagnoses data
diagnoses = pd.read_csv('DIAGNOSES_ICD.csv')
```

FILTERING FOR DIABETES DIAGNOSES

Title: Identifying Diabetes Patients

- Explanation: Filter the diagnoses to include only those with ICD-9 codes corresponding to Diabetes Mellitus.

```
# Load diagnoses data
diagnoses = pd.read_csv('DIAGNOSES_ICD.csv')

# Filter for Diabetes Mellitus ICD-9 codes
diabetes_codes = [f'250{str(i).zfill(2)}' for i in range(0, 100)]
diabetes_diagnoses = diagnoses[diagnoses['ICD9_CODE'].isin(diabetes_codes)].copy()

# Create a unique identifier for SUBJECT_ID and HADM_ID
diabetes_diagnoses.loc[:, 'SUBJ_HADM_ID'] = diabetes_diagnoses['SUBJECT_ID'].astype(str) + '_' +
diabetes_diagnoses['HADM_ID'].astype(str)

# Similarly, prepare the notes data
data['SUBJ_HADM_ID'] = data['SUBJECT_ID'].astype(str) + '_' + data['HADM_ID'].astype(str)

# Merge the notes with the Diabetes diagnoses
diabetes_notes = data[data['SUBJ_HADM_ID'].isin(diabetes_diagnoses['SUBJ_HADM_ID'])].copy()

# Reset index
diabetes_notes = diabetes_notes.reset_index(drop=True)
```


MERGING NOTES WITH DIAGNOSES

Title: Merging Clinical Notes with Diabetes Diagnoses

- **Explanation:** Create a unique identifier by combining SUBJECT_ID and HADM_ID, then merge the notes with the diabetes diagnoses based on this identifier.

```
# Create a unique identifier for SUBJECT_ID and HADM_ID
diabetes_diagnoses['SUBJ_HADM_ID'] = diabetes_diagnoses['SUBJECT_ID'].astype(str) + '_' + diabetes_diagnoses['HADM_ID'].astype(str)

# Similarly, prepare the notes data
data['SUBJ_HADM_ID'] = data['SUBJECT_ID'].astype(str) + '_' + data['HADM_ID'].astype(str)

# Merge the notes with the Diabetes diagnoses
diabetes_notes = data[data['SUBJ_HADM_ID'].isin(diabetes_diagnoses['SUBJ_HADM_ID'])]

# Reset index
diabetes_notes = diabetes_notes.reset_index(drop=True)
```

DATA CLEANING

Title: Cleaning Clinical Notes Text

- **Explanation:** Remove null entries in the TEXT column and clean the text by removing de-identification brackets and extra whitespace.

```
import re

# Remove null TEXT entries
diabetes_notes = diabetes_notes.dropna(subset=['TEXT'])

# Function to clean text
def clean_text(text):
    text = re.sub(r'\\[*\*.*?\*\\*\\]', '', text) # Remove de-identification brackets
    text = re.sub(r'\s+', ' ', text) # Remove extra whitespace
    return text.strip()

diabetes_notes['CLEAN_TEXT'] = diabetes_notes['TEXT'].apply(clean_text)
```

LOADING NLP MODELS

Title: Loading SpaCy and SciSpaCy Models

- **Explanation:** Load the necessary SpaCy and SciSpaCy models for NLP processing, including the biomedical NER model.

```
import spacy

# Load SpaCy's English model
nlp_spacy = spacy.load('en_core_web_sm')

# Load SciSpaCy's models
nlp_scispacy = spacy.load('en_core_sci_sm')

# Load SciSpaCy's NER model for biomedical data
nlp_scispacy_ner = spacy.load('en_ner_bc5cdr_md')
```


ENTITY RECOGNITION WITH SPACY

Title: Extracting Entities with SpaCy

- **Explanation:** Use SpaCy's English model to extract entities from a sample clinical note.

```
doc_spacy =  
nlp_spacy(diabetes_notes['CLEAN_TEXT'][0])  
print("SpaCy Entities:")  
for ent in doc_spacy.ents:  
    print(f"{ent.text} - {ent.label_}")
```

Output:

List of entities identified by SpaCy with their labels (e.g., Zocor / Lescol Attending - ORG).

```
SpaCy Entities:  
Zocor / Lescol Attending - ORG  
History of Present Illness -  
WORK_OF_ART  
84 - CARDINAL  
1 - CARDINAL  
28 - CARDINAL  
25-30% - PERCENT  
CAD - ORG  
s/p CABG - ORG  
SVG - ORG  
LIMA - GPE  
SVG-OM - ORG  
...
```


ENTITY RECOGNITION WITH SCISPACY

Title: Extracting Entities with SciSpaCy

- **Explanation:** Use SciSpaCy's model to extract scientific and biomedical entities from the same clinical note.

```
doc_scispacy =  
nlp_scispacy(diabetes_notes['CLEAN_TEXT'][0])  
print("\nSciSpaCy Entities:") for ent in  
doc_scispacy.ents: print(f"{ent.text} -  
{ent.label_}")
```

Output:

- List of entities identified by SciSpaCy with the label ENTITY.

```
SciSpaCy Entities:  
Admission - ENTITY  
Discharge Date - ENTITY  
Service - ENTITY  
MEDICINE Allergies - ENTITY  
Zocor - ENTITY  
Lescol - ENTITY  
Attending - ENTITY  
Chief Complaint - ENTITY  
Chest pain - ENTITY  
Surgical - ENTITY  
Invasive Procedure - ENTITY  
Central venous line insertion - ENTITY  
right internal jugular vein) History of Present -  
ENTITY  
Illness - ENTITY  
Mr. - ENTITY  
man - ENTITY  
...
```

BIOMEDICAL NER WITH SCISPACY

Title: Extracting Biomedical Entities with SciSpaCy NER

- **Explanation:** Use SciSpaCy's biomedical NER model to extract chemical and disease entities.

```
doc_scispacy_ner =  
nlp_scispacy_ner(diabetes_notes['CLEAN_TEXT'][0]  
)  
print("\nSciSpaCy NER Entities:")  
for ent in doc_scispacy_ner.ents:  
    print(f"{ent.text} - {ent.label_}")
```

Output:

- List of entities with labels **CHEMICAL** and **DISEASE** (e.g., Zocor - **CHEMICAL**, chest pain - **DISEASE**).

```
SciSpaCy NER Entities:  
Allergies - DISEASE  
Zocor - CHEMICAL  
Lescol - CHEMICAL  
Chest pain - DISEASE  
stenosis - DISEASE  
mitral regurgitation - DISEASE  
aortic insufficiency - DISEASE  
chronic left ventricular systolic heart failure - DISEASE  
hypertension - DISEASE  
hyperlipidemia - DISEASE  
diabetes mellitus - DISEASE  
CAD - DISEASE  
SVG-OM - CHEMICAL  
SVG-OM - CHEMICAL  
peripheral arterial disease - DISEASE  
...
```


VISUALIZING ENTITIES

Title: Visualizing Entities with displaCy

- **Explanation:** Use SpaCy's displaCy visualizer to display entities in the text.

Visualization:

- Annotated text highlighting entities such as medications, diseases, and procedures.

```
from spacy import displacy

# SpaCy Visualization
displacy.render(doc_spacy, style="ent", jupyter=True)

# SciSpaCy Visualization
displacy.render(doc_scispacy_ner, style="ent", jupyter=True)
```

Admission Date: Discharge Date: Service: MEDICINE Allergies: **Zocor / Lescor Attending ORG** : Chief Complaint: Chest pain Major Surgical or Invasive Procedure: Central venous line insertion (right internal jugular vein) **History of Present Illness WORK_OF_ART** : Mr. is an **84 CARDINAL** yo man with moderate aortic stenosis (outside hospital echo in with **1 CARDINAL** cm2, gradient **28 CARDINAL** mmHg, moderate mitral regurgitation, mild aortic insufficiency), chronic left ventricular systolic heart failure with EF **25-30% PERCENT** , hypertension, hyperlipidemia, diabetes mellitus, **CAD ORG** **s/p CABG ORG** in with SVG-LAD-Diagonal, **SVG ORG** -OM, and SVG-RPDA-RPL, with a re-do CABG in with **LIMA GPE** -LAD, **SVG-OM ORG** , **SVG ORG** -diagonal, and SVG-RCA. He also has severe peripheral arterial disease s/p peripheral bypass surgery. He presented to **Hospital ER FAC** this morning **TIME** with shortness of breath and chest pain and was found to be in heart failure. He states he was in his usual state of health until **10:30 last evening TIME** when he woke up feeling cold; **1 hour TIME** later he developed moderate to severe sharp chest pain radiating across his chest associated with nausea, diaphoresis, and dyspnea. The pain was fairly constant and did not resolve until he was given sL NTG at **6 am TIME** by

PREPROCESSING TEXT FOR WORD2VEC

Title: Preprocessing Text Data

- **Explanation:** Tokenize and preprocess the cleaned text for Word2Vec modeling, removing stopwords and short tokens.

```
from gensim.utils import simple_preprocess
from gensim.parsing.preprocessing import STOPWORDS

def preprocess(text):
    return [token for token in simple_preprocess(text) if token not in STOPWORDS and len(token) > 3]

documents = diabetes_notes['CLEAN_TEXT'].map(preprocess)
```


TRAINING WORD2VEC MODEL

Title: Training a Word2Vec Model on Clinical Notes

- **Explanation:** Train a Word2Vec model using the preprocessed documents to learn word embeddings.

```
from gensim.models import Word2Vec

model = Word2Vec(sentences=documents,
vector_size=100, window=5, min_count=2, workers=4)

# Most similar words to 'insulin'
print("Words similar to 'insulin':")
print(model.wv.most_similar('insulin'))
```

Output:

```
Words similar to 'insulin':
[
('humalog', 0.7126625180244446),
('lantus', 0.6957399845123291),
('insuling', 0.6520899534225464),
('hiss', 0.6348435878753662),
('insuln', 0.6166067719459534),
('riss', 0.595574140548706),
('fingersticks', 0.5858661532402039),
('glargin', 0.5817741751670837),
('units', 0.555779755115509),
('ultralente', 0.5463814735412598)
]
```

VISUALIZING WORD EMBEDDINGS WITH T-SNE

Title: Visualizing Word2Vec Embeddings

- **Explanation:** Use t-SNE to reduce the dimensionality of word embeddings and visualize them.

```
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import numpy as np

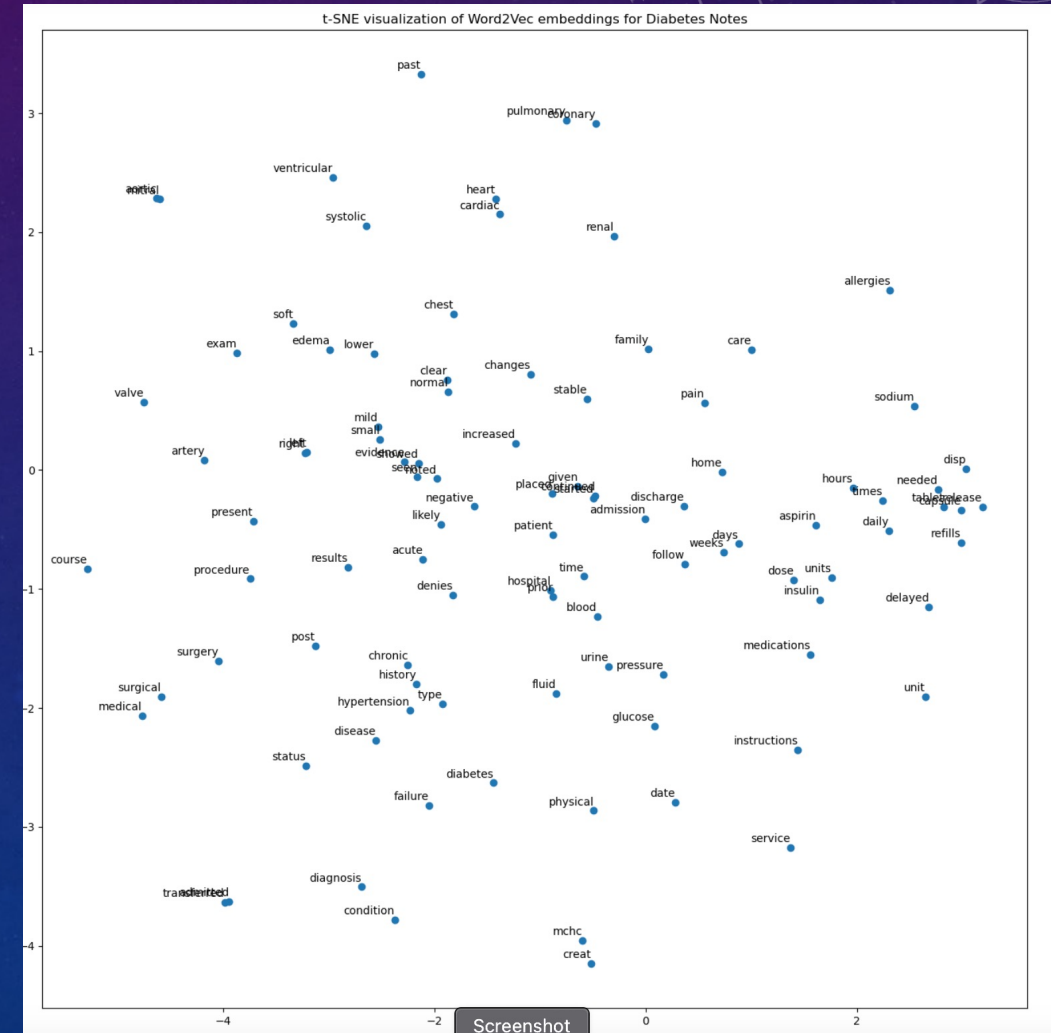
# Get vocabulary and limit to top N words
vocab = list(model.wv.key_to_index)[:100]

# Get word vectors
X = model.wv[vocab]

tsne = TSNE(n_components=2, perplexity=30, random_state=42)
X_tsne = tsne.fit_transform(X)

plt.figure(figsize=(16, 16))
plt.scatter(X_tsne[:, 0], X_tsne[:, 1])

for i, word in enumerate(vocab):
    plt.annotate(word, xy=(X_tsne[i, 0], X_tsne[i, 1]), textcoords='offset
points', xytext=(0, 5), ha='right')
plt.title('t-SNE Visualization of Word2Vec Embeddings for Diabetes Notes')
plt.show()
```



USING CLINICALBERT FOR EMBEDDINGS

Title: Generating Embeddings with ClinicalBERT

- **Explanation:** Use the emilyalsentzer/Bio_ClinicalBERT model to generate embeddings for a clinical note.

```
from transformers import AutoTokenizer, AutoModel
import torch

tokenizer = AutoTokenizer.from_pretrained("emilyalsentzer/Bio_ClinicalBERT")
model_bert = AutoModel.from_pretrained("emilyalsentzer/Bio_ClinicalBERT")

# Tokenize and encode a sample note
inputs = tokenizer(diabetes_notes['CLEAN_TEXT'][0], return_tensors="pt", max_length=512, truncation=True)

# Get the embeddings
with torch.no_grad():
    outputs = model_bert(**inputs)

embeddings = outputs.last_hidden_state # Shape: [batch_size, sequence_length, hidden_size]
```

VISUALIZING CLINICALBERT EMBEDDINGS

Title: Visualizing Token Embeddings with t-SNE

- **Explanation:** Reduce the dimensions of the token embeddings and visualize them using t-SNE.

```
# Convert embeddings to numpy for visualization
token_embeddings = embeddings[0].numpy()


# Get tokens
tokens = tokenizer.tokenize(diabetes_notes['CLEAN_TEXT'][0], max_length=512, truncation=True)

# Reduce dimensions
tsne = TSNE(n_components=2, perplexity=30, random_state=42)
embeddings_tsne = tsne.fit_transform(token_embeddings)

# Plot
plt.figure(figsize=(16, 16))
plt.scatter(embeddings_tsne[:, 0], embeddings_tsne[:, 1])

for i, token in enumerate(tokens):
    plt.annotate(token, xy=(embeddings_tsne[i, 0], embeddings_tsne[i, 1]), textcoords='offset points', xytext=(0, 5),
ha='right')
plt.title('t-SNE Visualization of ClinicalBERT Embeddings for Diabetes Notes')
plt.show()
```


t-SNE visualization of ClinicalBERT embeddings for Diabetes Notes



COMPARING CLINICALBERT WITH STANDARD BERT

Title: Tokenization Comparison between ClinicalBERT and BERT

- **Explanation:** Compare the tokens generated by ClinicalBERT and standard BERT models to highlight differences in handling clinical text.

```
# Load standard BERT tokenizer
from transformers import BertTokenizer, BertModel

bert_tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
bert_model = BertModel.from_pretrained('bert-base-uncased')

# Tokenize with standard BERT
bert_tokens = bert_tokenizer.tokenize(diabetes_notes['CLEAN_TEXT'][0], max_length=512, truncation=True)

# Compare tokens
print("Standard BERT Tokens:")
print(bert_tokens)

print("\nClinicalBERT Tokens:")
print(tokens)
```


Output:

- Lists of tokens from both models showing that ClinicalBERT better handles medical terminology.

Standard BERT Tokens:

```
['admission', 'date', ':', 'discharge', 'date', ':', 'service', ':', 'medicine', 'all', '##er',  
##gies', ':', 'z', '##oco', '##r', '/', 'les', '##col', 'attending', ':', 'chief', 'complaint',  
:', 'chest', 'pain', 'major', 'surgical', 'or', 'invasive', 'procedure', ':', 'central', 've',  
##nous', 'line', 'insertion', '(', 'right', 'internal', 'jug', '##ular', 'vein', ')', 'history',  
of', 'present', 'illness', ':', 'mr', '.', 'is', 'an', '84', 'yo', 'man', 'with', 'moderate',  
ao', '##rti', '##c', 'ste', '##nosis', '(', 'outside', 'hospital', 'echo', 'in', 'with', '1',  
cm', '##2', ',', 'gradient', '28', 'mm', '##hg', ',', 'moderate', 'mit', '##ral', 'reg',  
##urg', '##itation', ',', 'mild', 'ao', '##rti', '##c', 'ins', '##uf', '##fi', '##ciency', ')',  
,', 'chronic', 'left', 'vent', '##ric', '##ular', 'sy', '##sto', '##lic', 'heart', 'failure',  
with', 'e', '##f', '25', '-', '30', '%', ',', 'hyper', '##tension', ',', 'hyper', '##lip',  
##ide', '##mia', ',', 'diabetes', 'mel', '##lit', '##us', ',', 'cad', 's', '/', 'p', 'cab',  
##g', 'in', 'with', 'sv', '##g', '-', 'lad', '-', 'diagonal', ',', 'sv', '##g', '-', 'om', ',',  
and', 'sv', '##g', '-', 'r', '##pd', '##a', '-', 'r', '##pl', ',', 'with', 'a', 're', '-', 'do',  
cab', '##g', 'in', 'with', 'lima', '-', 'lad', ',', 'sv', '##g', '-', 'om', ',', 'sv', '##g', '-',  
, 'diagonal', ',', 'and', 'sv', '##g', '-', 'rca', '.', 'he', 'also', 'has', 'severe',  
'peripheral', 'arterial', 'disease', 's', '/', 'p', 'peripheral', 'bypass', 'surgery', '.', 'he',  
'presented', 'to', 'hospital', 'er', 'this', 'morning', 'with', 'short', ...]
```

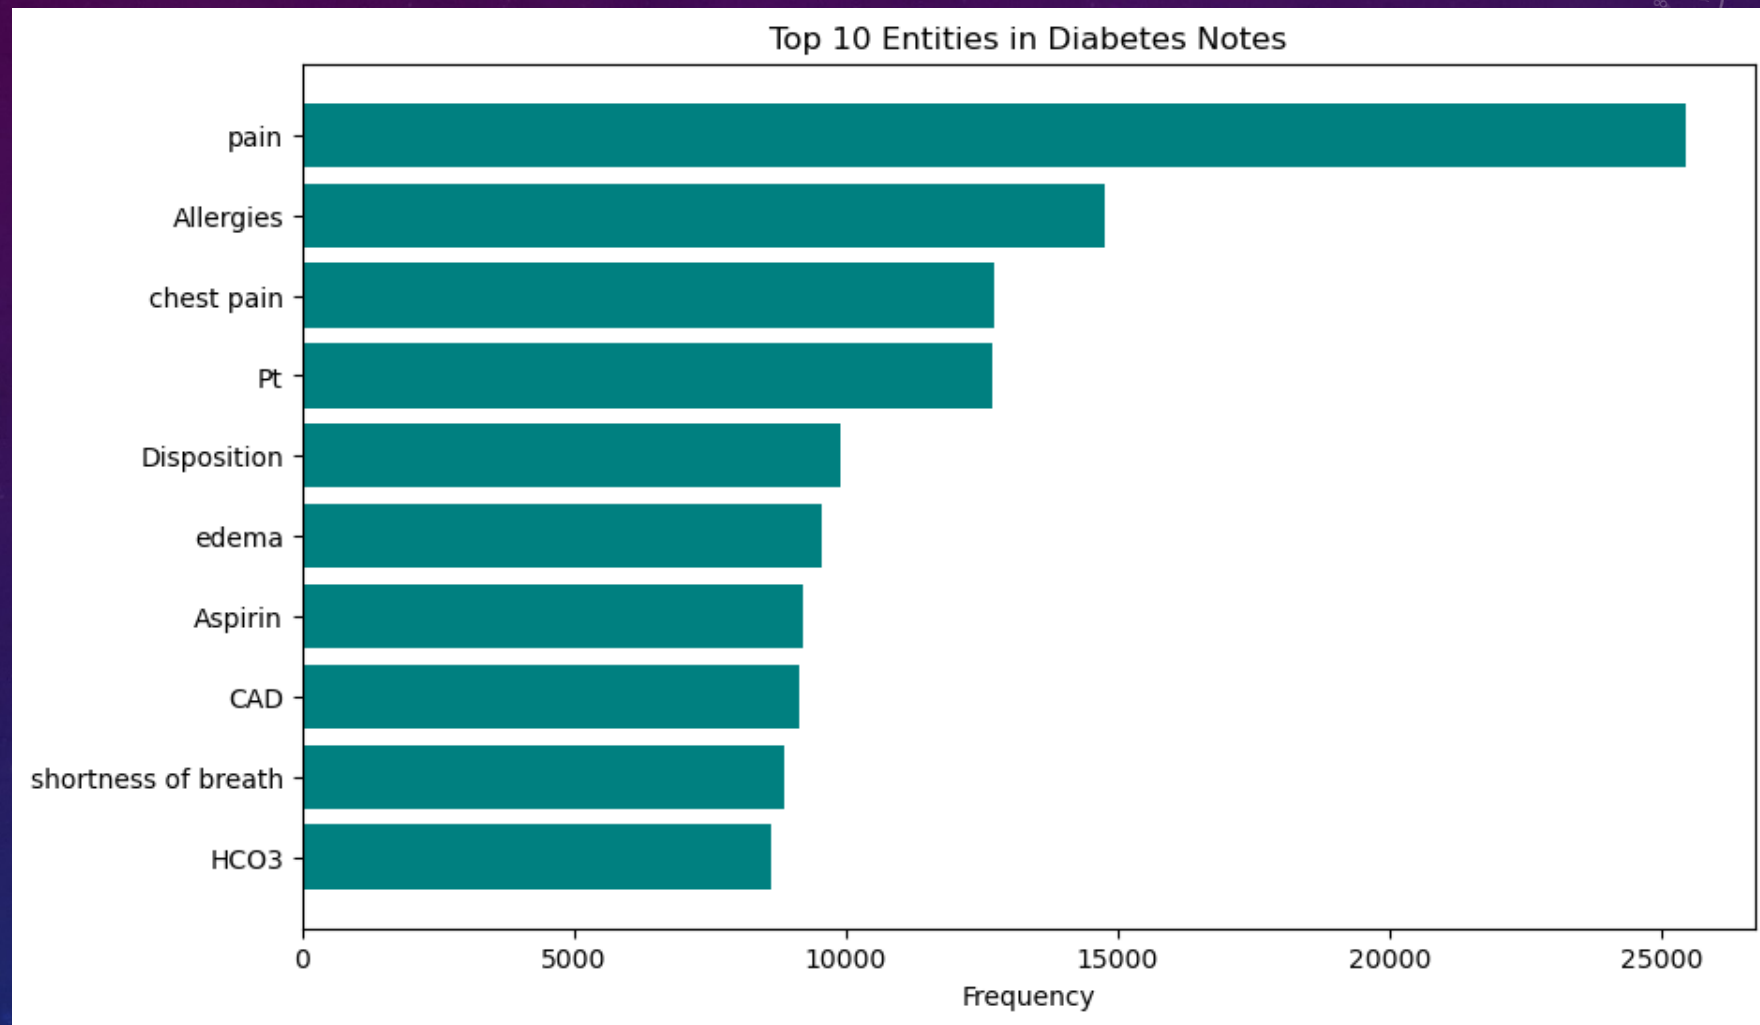
TOP 10 ENTITIES FROM ALL DIABETES NOTES

```
# Extract entities from all diabetes notes
entities = []
for text in diabetes_notes['CLEAN_TEXT']:
    doc = nlp_scispacy_ner(text)
    entities.extend([ent.text for ent in doc.ents])

# Count entity frequencies
from collections import Counter
entity_counts = Counter(entities)
most_common_entities = entity_counts.most_common(10)

# Separate labels and counts
labels, counts = zip(*most_common_entities)

# Plot
plt.figure(figsize=(10, 6))
plt.barh(labels, counts, color='teal')
plt.gca().invert_yaxis()
plt.xlabel('Frequency')
plt.title('Top 10 Entities in Diabetes Notes')
plt.show()
```

OBSERVATIONS AND INSIGHTS

Title: Key Takeaways from NLP Analysis

- **Explanation:**
 - **Entity Recognition:**
 - SpaCy provides general-purpose entity recognition.
 - SciSpaCy and its NER model effectively identify biomedical entities like diseases and chemicals.
 - **Word Embeddings:**
 - Word2Vec captures semantic relationships between medical terms (e.g., 'insulin' is similar to 'lantus' and 'humalog').
 - ClinicalBERT generates contextual embeddings tailored for clinical text.
 - **Tokenization Differences:**
 - ClinicalBERT tokenizes medical terms more appropriately than standard BERT, leading to better representations.